

Summer boot camp project 2024

YASH CHAUHAN

List Of Tables

List of figures

Problem Statement

The data contains the details for the Uber rides across various boroughs (subdivisions) of New York City at an hourly level and attributes associated with weather conditions at that time.

Importing the Important libraries

```
In [1]:  ▶ import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

Loading the Dataset

```
In [2]:  ▶ df = pd.read_csv("D:/dataset.csv")
```

1- Display first five rows of teh dataset

In [3]: `df.head()`

Out[3]:

	pickup_dt	borough	pickups	spd	vsb	temp	dewp	slp	pcp01	pcp06	pcp24
0	01-01-2015 01:00	Bronx	152.0	5.0	10.0	30.0	7.0	1023.5	0.0	0.0	0.0
1	01-01-2015 01:00	Brooklyn	1519.0	5.0	10.0	NaN	7.0	1023.5	0.0	0.0	0.0
2	01-01-2015 01:00	EWB	0.0	5.0	10.0	30.0	7.0	1023.5	0.0	0.0	0.0
3	01-01-2015 01:00	Manhattan	5258.0	5.0	10.0	30.0	7.0	1023.5	0.0	0.0	0.0
4	01-01-2015 01:00	Queens	405.0	5.0	10.0	30.0	7.0	1023.5	0.0	0.0	0.0

	pickup_dt	borough	pickups	spd	vsb	temp	dewp	slp	pcp01	pcp06	pcp24	sd	hday
0	1/1/2015 1:00	Bronx	152.0	5.0	10.0	30.0	7.0	1023.5	0.0	0.0	0.0	0.0	Y
1	1/1/2015 1:00	Brooklyn	1519.0	5.0	10.0	NaN	7.0	1023.5	0.0	0.0	0.0	0.0	?
2	1/1/2015 1:00	EWB	0.0	5.0	10.0	30.0	7.0	1023.5	0.0	0.0	0.0	0.0	Y
3	1/1/2015 1:00	Manhattan	5258.0	5.0	10.0	30.0	7.0	1023.5	0.0	0.0	0.0	0.0	Y
4	1/1/2015 1:00	Queens	405.0	5.0	10.0	30.0	7.0	1023.5	0.0	0.0	0.0	0.0	Y

Observations:

- First five rows of the dataset
- In first five rows we can see a "?" in the hday values.

2- Display last five rows of the dataset

In [4]: `df.tail()`

Out[4]:

	pickup_dt	borough	pickups	spd	vsb	temp	dewp	slp	pcp01	pcp06	p
29096	30-06-2015 23:00	EWB	0.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	
29097	30-06-2015 23:00	Manhattan	3828.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	
29098	30-06-2015 23:00	Queens	580.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	
29099	30-06-2015 23:00	Staten Island	0.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	
29100	30-06-2015 23:00	NaN	3.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	

	pickup_dt	borough	pickups	spd	vsb	temp	dewp	slp	pcp01	pcp06	pcp24	sd	hday
29096	30-06-2015 23:00	EWB	0.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	0.0	0.0	N
29097	30-06-2015 23:00	Manhattan	3828.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	0.0	0.0	N
29098	30-06-2015 23:00	Queens	580.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	0.0	0.0	N
29099	30-06-2015 23:00	Staten Island	0.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	0.0	0.0	N
29100	30-06-2015 23:00	NaN	3.0	7.0	10.0	75.0	65.0	1011.8	0.0	0.0	0.0	0.0	N

3- Check for shape of the dataset

In [5]: `df.shape`

Out[5]: (29101, 13)

(29101, 13)

Observations:

- There are 29101 number of rows in our dataset and 13 columns

4- Check the datatypes of each feature.

In [6]: `df.dtypes`

```
Out[6]: pickup_dt      object
borough      object
pickups      float64
spd          float64
vsb          float64
temp         float64
dewp         float64
slp          float64
pcp01        float64
pcp06        float64
pcp24        float64
sd           float64
hday         object
dtype: object
```

```
pickup_dt      object
borough      object
pickups      float64
spd          float64
vsb          float64
temp         float64
dewp         float64
slp          float64
pcp01        float64
pcp06        float64
pcp24        float64
sd           float64
hday         object
dtype: object
```

5- Check the Statistical summary

In [7]: `df.describe()`

```
Out[7]:
```

	pickups	spd	vsb	temp	dewp	
count	29099.000000	29101.000000	29101.000000	28742.000000	29101.000000	2.910100e
mean	490.236022	5.984924	8.818125	47.900262	30.823065	1.052633e
std	995.680628	3.699007	2.442897	19.800541	21.283444	5.945147e
min	0.000000	0.000000	0.000000	0.000000	-16.000000	1.000000e
25%	1.000000	3.000000	9.100000	32.000000	14.000000	1.012500e
50%	54.000000	6.000000	10.000000	46.500000	30.000000	1.018200e
75%	449.000000	8.000000	10.000000	65.000000	50.000000	1.022900e
max	7883.000000	21.000000	10.000000	89.000000	73.000000	1.015200e

	pickups	spd	vsb	temp	dewp	slp	pcp01	pcp06	pcp24	sd
count	29099.000000	29101.000000	29101.000000	28742.000000	29101.000000	2.910100e+04	29101.000000	29101.000000	29101.000000	29101.000000
mean	490.236022	5.984924	8.818125	47.900262	30.823065	1.052633e+03	0.003830	0.026129	0.090464	2.529169
std	995.680628	3.699007	2.442897	19.800541	21.283444	5.945147e+03	0.018933	0.093125	0.219402	4.520325
min	0.000000	0.000000	0.000000	0.000000	-16.000000	1.000000e+00	0.000000	0.000000	0.000000	0.000000
25%	1.000000	3.000000	9.100000	32.000000	14.000000	1.012500e+03	0.000000	0.000000	0.000000	0.000000
50%	54.000000	6.000000	10.000000	46.500000	30.000000	1.018200e+03	0.000000	0.000000	0.000000	0.000000
75%	449.000000	8.000000	10.000000	65.000000	50.000000	1.022900e+03	0.000000	0.000000	0.050000	2.958333
max	7883.000000	21.000000	10.000000	89.000000	73.000000	1.015200e+06	0.280000	1.240000	2.100000	19.000000

6- Check the null values

In [8]: `df.isnull().sum()`

```
Out[8]: pickup_dt      0
        borough    3043
        pickups      2
        spd         0
        vsb         0
        temp      359
        dewp        0
        slp         0
        pcp01        0
        pcp06        0
        pcp24        0
        sd           0
        hday         0
        dtype: int64
```

```
        pickup_dt      0
        borough    3043
        pickups      2
        spd         0
        vsb         0
        temp      359
        dewp        0
        slp         0
        pcp01        0
        pcp06        0
        pcp24        0
        sd           0
        hday         0
        dtype: int64
```

Observations:

- There are 3043 null values in our dataset in borough column.
- There are 2 null values in pickup column and 359 null values in the temp column of the dataset.

7- Check the duplicate values

```
In [9]: ▶ duplicates = df.duplicated().sum()
print(f"Number of duplicate rows: {duplicates}")
```

Number of duplicate rows: 0

Number of duplicate rows: 0

8- Check for outliers and their authenticity

```
In [10]: ▶ numerical_columns = df.select_dtypes(include=['float64']).columns
# Calculate IQR to identify outliers
Q1 = df[numerical_columns].quantile(0.25)
Q3 = df[numerical_columns].quantile(0.75)
IQR = Q3 - Q1

# Identify outliers
outliers = (df[numerical_columns] < (Q1 - 1.5 * IQR)) | (df[numerical_c

# Display outliers summary
outliers_summary = outliers.sum()
print("Number of outliers in each column:")
print(outliers_summary)
```

Number of outliers in each column:

```
pickups    3498
spd         451
vsb        5322
temp         0
dewp         0
slp         281
pcp01       2633
pcp06       5641
pcp24       5016
sd          6060
dtype: int64
```

Number of outliers in each column:

```
pickups    3498
spd         451
vsb        5322
temp         0
dewp         0
slp         281
pcp01       2633
pcp06       5641
pcp24       5016
sd          6060
dtype: int64
```

```
In [11]: ► outlier_counts = outliers.sum()
total_counts = df.shape[0]
outlier_percentage = (outlier_counts / total_counts) * 100

# Display the percentage of outliers
print("Percentage of outliers in each column:")
print(outlier_percentage)
```

Percentage of outliers in each column:

pickups	12.020205
spd	1.549775
vsb	18.288031
temp	0.000000
dewp	0.000000
slp	0.965603
pcp01	9.047799
pcp06	19.384214
pcp24	17.236521
sd	20.824027

dtype: float64

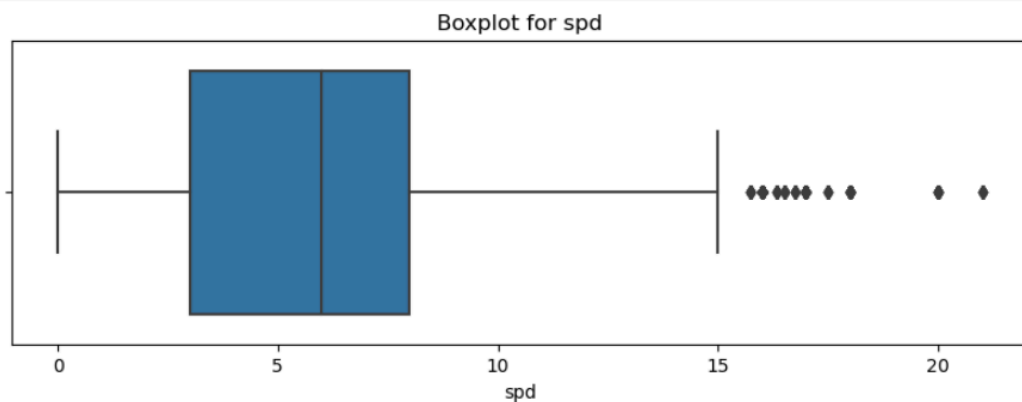
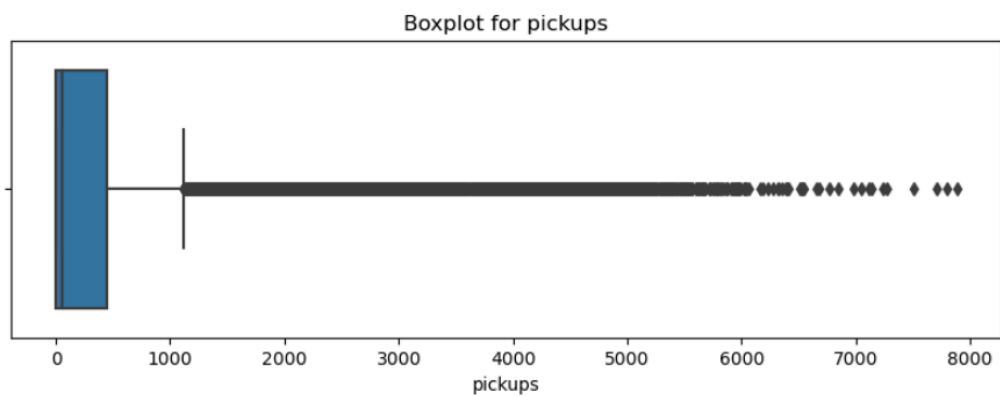

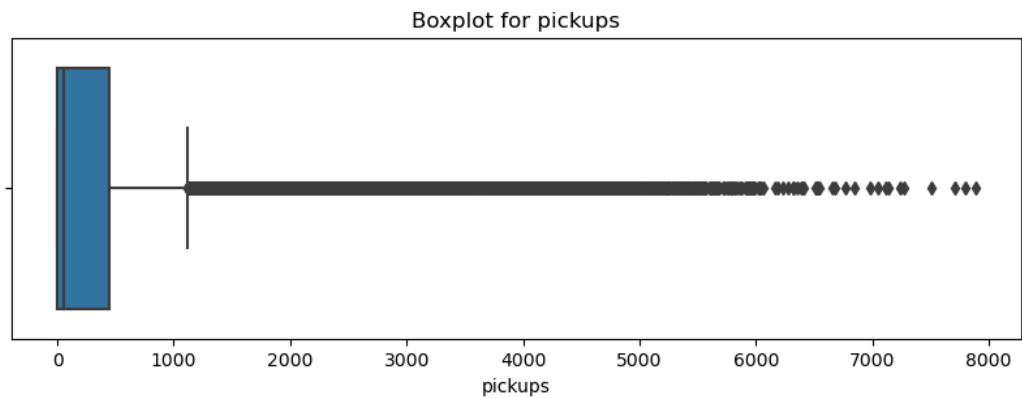
Percentage of outliers in each column:

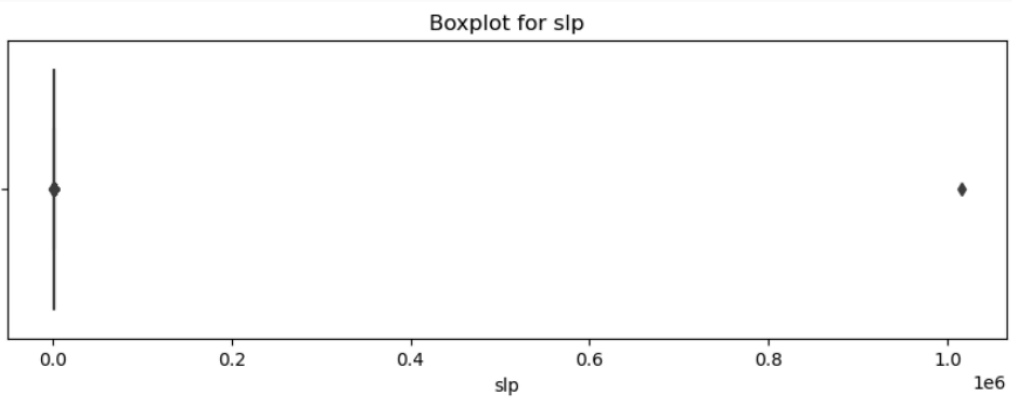
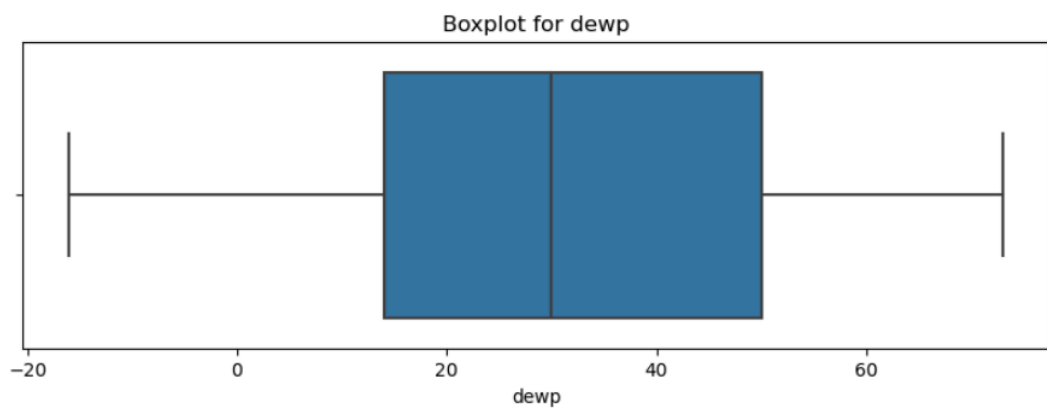

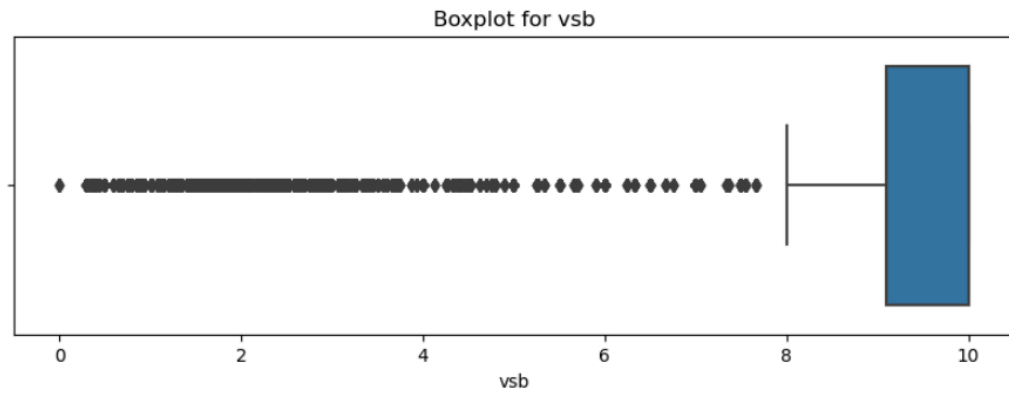
pickups	12.020205
spd	1.549775
vsb	18.288031
temp	0.000000
dewp	0.000000
slp	0.965603
pcp01	9.047799
pcp06	19.384214
pcp24	17.236521
sd	20.824027

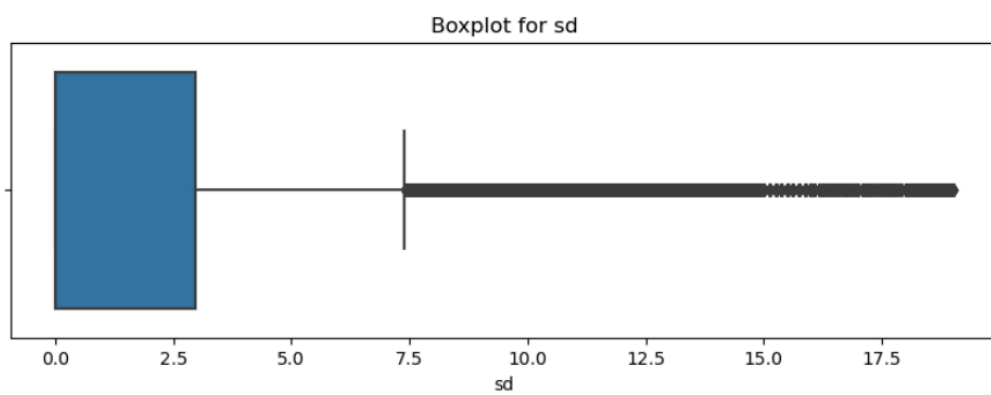
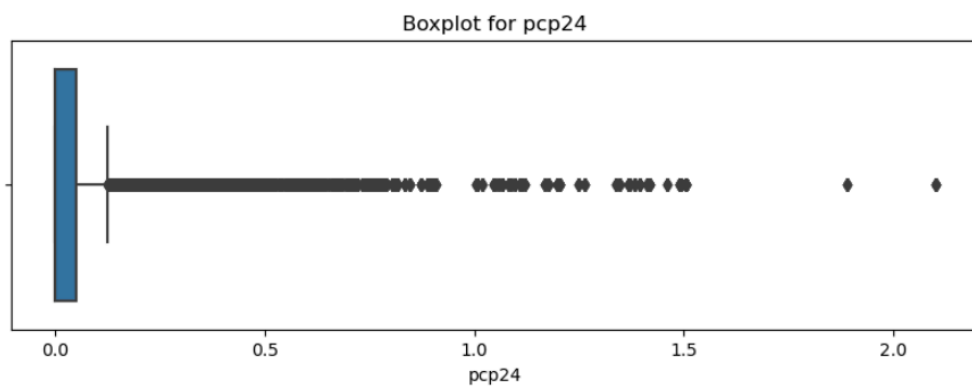

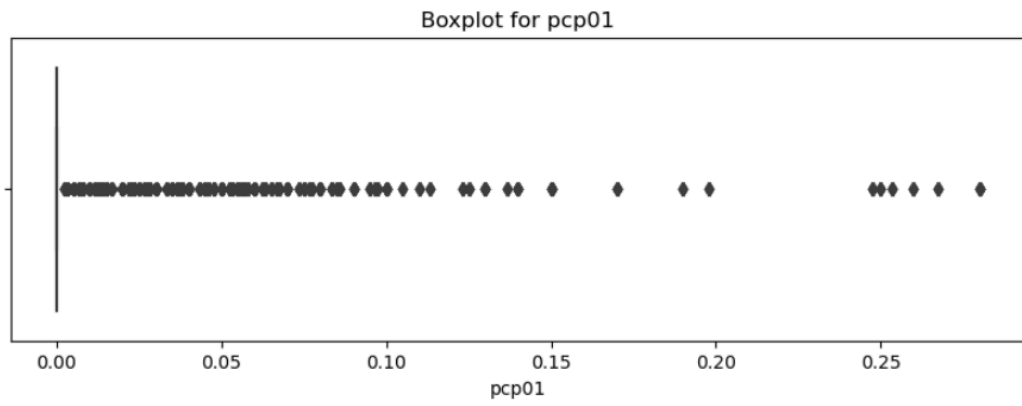
dtype: float64

```
In [12]: numerical_columns = df.select_dtypes(include=['float64']).columns

for column in numerical_columns:
    plt.figure(figsize=(10, 3))
    sns.boxplot(x=df[column])
    plt.title(f'Boxplot for {column}')
    plt.show()
```







9- Check for Anomalies:

In [13]:

```

columns_with_question_mark = {}
for column in df.columns:
    count = (df[column] == '?').sum()
    if count > 0:
        columns_with_question_mark[column] = count

print("Columns with '?' and their counts:")
for column, count in columns_with_question_mark.items():
    print(f"{column}: {count}")

```

Columns with '?' and their counts:
hday: 2

Columns with '?' and their counts:
hday: 2

In [14]:

```

# Filter rows where 'hday' column has '?'
rows_with_question_mark= df[df['hday'] == '?']

print("Rows with '?' in 'hday' column:")
print(rows_with_question_mark)

```

Rows with '?' in 'hday' column:

		pickup_dt	borough	pickups	spd	vsb	temp	dewp	sl
p \									
1	01-01-2015	01:00	Brooklyn	1519.0	5.0	10.0	NaN	7.0	1023.5
123	01-01-2015	19:00	Queens	238.0	7.0	10.0	37.0	7.0	1016.2

	pcp01	pcp06	pcp24	sd	hday
1	0.0	0.0	0.0	0.0	?
123	0.0	0.0	0.0	0.0	?

Rows with '?' in 'hday' column:

	pickup_dt	borough	pickups	spd	vsb	temp	dewp	slp	pcp01	\
1	1/1/2015 1:00	Brooklyn	1519.0	5.0	10.0	NaN	7.0	1023.5	0.0	
123	1/1/2015 19:00	Queens	238.0	7.0	10.0	37.0	7.0	1016.2	0.0	

	pcp06	pcp24	sd	hday
1	0.0	0.0	0.0	?
123	0.0	0.0	0.0	?

Observations:

- We have got dtypes of date as object that should be date and time.
- There are two rows which have "?" as the value in their hday value.
- The number of outliers is much more in our dataset percentage ranging from 12-20%.

10- Necessary cleaning needed

```
In [15]: df.replace('?', np.nan, inplace=True)
```

```
In [16]: rows_with_question_mark= df[df['hday'] == '?']
```

```
print("Rows with '?' in 'hday' column:")  
print(rows_with_question_mark)
```

Rows with '?' in 'hday' column:

Empty DataFrame

Columns: [pickup_dt, borough, pickups, spd, vsb, temp, dewp, slp, pcp01, pcp06, pcp24, sd, hday]

Index: []

Rows with '?' in 'hday' column:

Empty DataFrame

Columns: [pickup_dt, borough, pickups, spd, vsb, temp, dewp, slp, pcp01, pcp06, pcp24, sd, hday]

Index: []

Removing null values using mean, median and mode

```
In [17]: # Filling missing values in pickups column with the median of the column  
df['pickups'].replace(np.nan, df['pickups'].median(), inplace=True)
```

```
In [18]: # Filling null values in borough column with the mode of the column  
df['borough'].replace(np.nan, df['borough'].mode().values[0], inplace=True)
```

```
In [19]: # Filling missing values in temp columns with the median of the column  
df['temp'].replace(np.nan, df['temp'].median(), inplace=True)
```

```
In [20]: # Filling null values in hday column with the mode of the column  
df['hday'].replace(np.nan, df['hday'].mode().values[0], inplace=True)
```

```
In [21]: df.isnull().sum()
```

```
Out[21]: pickup_dt      0  
borough      0  
pickups      0  
spd          0  
vsb          0  
temp         0  
dewp         0  
slp          0  
pcp01        0  
pcp06        0  
pcp24        0  
sd           0  
hday         0  
dtype: int64
```

```
pickup_dt      0  
borough      0  
pickups      0  
spd          0  
vsb          0  
temp         0  
dewp         0  
slp          0  
pcp01        0  
pcp06        0  
pcp24        0  
sd           0  
hday         0  
dtype: int64
```

Removing Outliers from the dataset as data cleaning process

```
In [22]: #Function to remove the Outliers
```

```
def remove_outlier(col):  
    sorted(col)  
    Q1,Q3=col.quantile([0.25,0.75])  
    IQR=Q3-Q1  
    lower_range= Q1-(1.5 * IQR)  
    upper_range= Q3+(1.5 * IQR)  
    return lower_range, upper_range  
  
for i in df.columns:  
    if df[i].dtype !='object':  
        lr,ur=remove_outlier(df[i])  
        df[i]=np.where(df[i]>ur,ur,df[i])  
        df[i]=np.where(df[i]<lr,lr,df[i])
```

```
In [23]: ▶ # Calculate IQR to identify outliers
Q1 = df[numerical_columns].quantile(0.25)
Q3 = df[numerical_columns].quantile(0.75)
IQR = Q3 - Q1

# Identify outliers
outliers = (df[numerical_columns] < (Q1 - 1.5 * IQR)) | (df[numerical_c

# Display outliers summary
outliers_summary = outliers.sum()
print("Number of outliers in each column:")
print(outliers_summary)
```

Number of outliers in each column:

```
pickups    0
spd        0
vsb        0
temp       0
dewp       0
slp        0
pcp01      0
pcp06      0
pcp24      0
sd         0
dtype: int64
```

Number of outliers in each column:

```
pickups    0
spd        0
vsb        0
temp       0
dewp       0
slp        0
pcp01      0
pcp06      0
pcp24      0
sd         0
dtype: int64
```

Observations:

- No null values in the dataset.
- Now there is no outliers in our data.
- Since we have done most of the necessary cleaning of our data, now we will perform EDA on it.

In [24]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29101 entries, 0 to 29100
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   pickup_dt   29101 non-null  object
1   borough     29101 non-null  object
2   pickups     29101 non-null  float64
3   spd         29101 non-null  float64
4   vsb         29101 non-null  float64
5   temp        29101 non-null  float64
6   dewp        29101 non-null  float64
7   slp         29101 non-null  float64
8   pcp01       29101 non-null  float64
9   pcp06       29101 non-null  float64
10  pcp24       29101 non-null  float64
11  sd          29101 non-null  float64
12  hday        29101 non-null  object
dtypes: float64(10), object(3)
memory usage: 2.9+ MB
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29101 entries, 0 to 29100
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   pickup_dt   29101 non-null  object
1   borough     29101 non-null  object
2   pickups     29101 non-null  float64
3   spd         29101 non-null  float64
4   vsb         29101 non-null  float64
5   temp        29101 non-null  float64
6   dewp        29101 non-null  float64
7   slp         29101 non-null  float64
8   pcp01       29101 non-null  float64
9   pcp06       29101 non-null  float64
10  pcp24       29101 non-null  float64
11  sd          29101 non-null  float64
12  hday        29101 non-null  object
dtypes: float64(10), object(3)
memory usage: 2.9+ MB
```

In [25]: `#Statistical summary
df.describe()`

Out[25]:

	pickups	spd	vsb	temp	dewp	
count	29101.000000	29101.000000	29101.000000	29101.000000	29101.000000	29101.000000
mean	282.045256	5.961970	9.487576	47.882988	30.823065	1017.810000
std	386.660889	3.631521	0.906658	19.678631	21.283444	7.700000
min	0.000000	0.000000	7.750000	0.000000	-16.000000	996.900000
25%	1.000000	3.000000	9.100000	32.000000	14.000000	1012.500000
50%	54.000000	6.000000	10.000000	46.500000	30.000000	1018.200000
75%	449.000000	8.000000	10.000000	64.500000	50.000000	1022.900000
max	1121.000000	15.500000	10.000000	89.000000	73.000000	1038.500000

	pickups	spd	vsb	temp	dewp	slp	pcp01	pcp06	pcp24	sd
count	29101.000000	29101.000000	29101.000000	29101.000000	29101.000000	29101.000000	29101.0	29101.0	29101.000000	29101.000000
mean	282.045256	5.961970	9.487576	47.882988	30.823065	1017.810618	0.0	0.0	0.030223	1.861165
std	386.660889	3.631521	0.906658	19.678631	21.283444	7.701187	0.0	0.0	0.049091	3.104397
min	0.000000	0.000000	7.750000	0.000000	-16.000000	996.900000	0.0	0.0	0.000000	0.000000
25%	1.000000	3.000000	9.100000	32.000000	14.000000	1012.500000	0.0	0.0	0.000000	0.000000
50%	54.000000	6.000000	10.000000	46.500000	30.000000	1018.200000	0.0	0.0	0.000000	0.000000
75%	449.000000	8.000000	10.000000	64.500000	50.000000	1022.900000	0.0	0.0	0.050000	2.958333
max	1121.000000	15.500000	10.000000	89.000000	73.000000	1038.500000	0.0	0.0	0.125000	7.395833

1- Pickup Analysis

In [26]: `# 1. Total number of Uber pickups across all boroughs
total_pickups = df['pickups'].sum()
total_pickups`

Out[26]: 8207799.0

8207799.0


```
In [27]: ▶ avg_hourly_pickups = df.groupby('borough')['pickups'].mean()
highest_avg_borough = avg_hourly_pickups.idxmax()
highest_avg_pickups = avg_hourly_pickups.max()

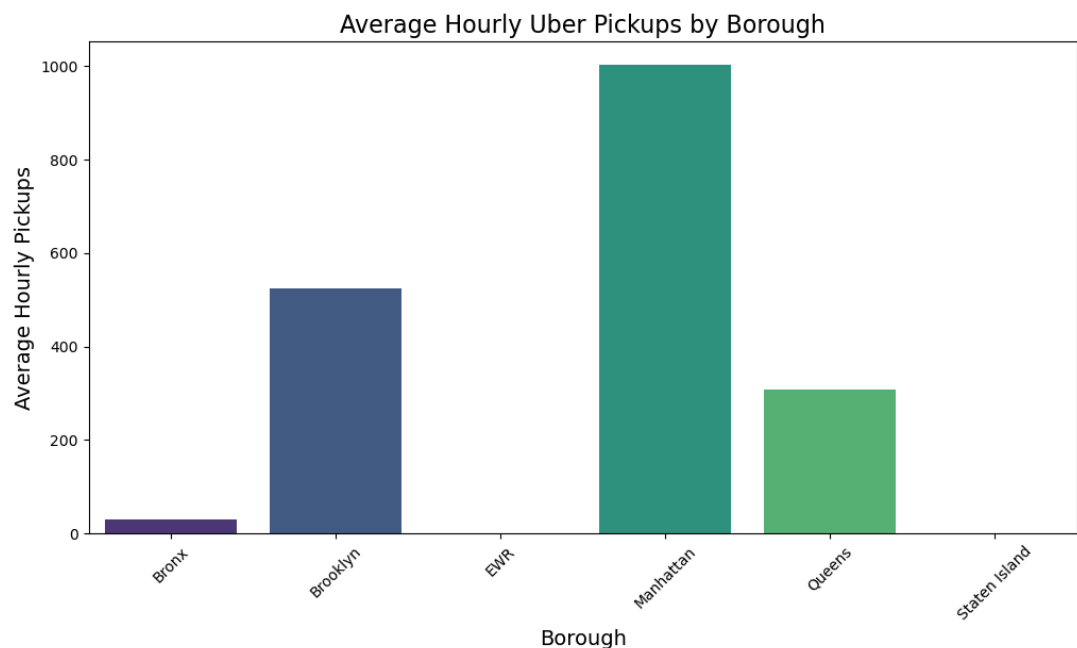
print(f"Borough with highest average hourly pickups: {highest_avg_borough}")

# Plotting the data
plt.figure(figsize=(12, 6))
sns.barplot(x=avg_hourly_pickups.index, y=avg_hourly_pickups.values, palette='magma')

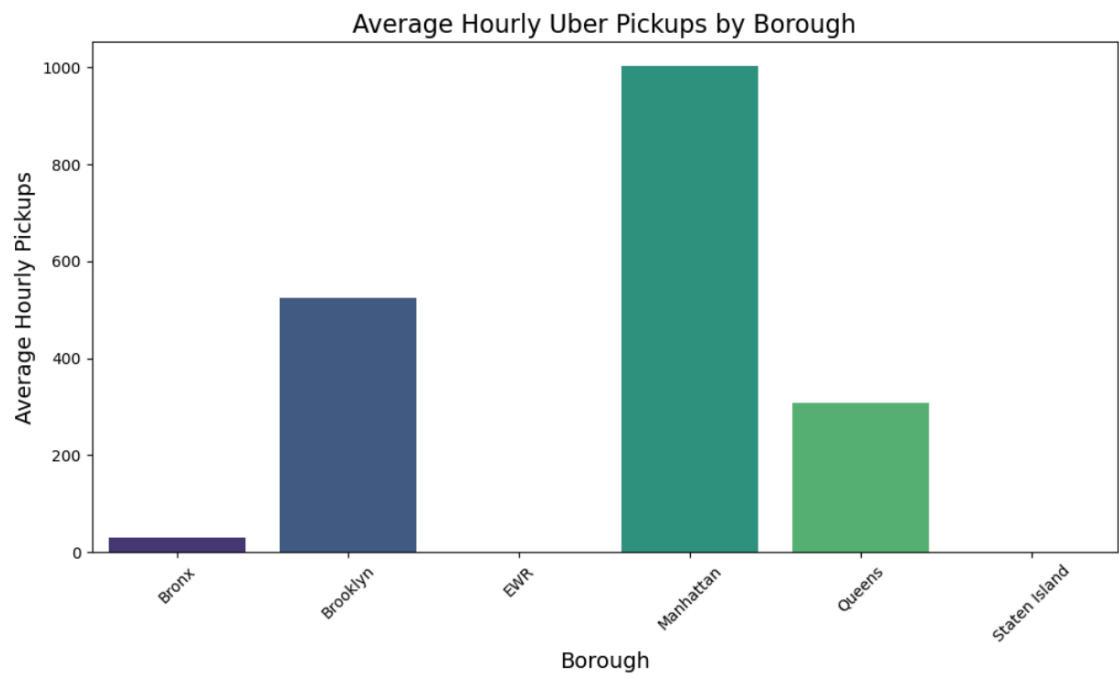
plt.title('Average Hourly Uber Pickups by Borough', fontsize=16)
plt.xlabel('Borough', fontsize=14)
plt.ylabel('Average Hourly Pickups', fontsize=14)
plt.xticks(rotation=45)


# Display the plot
plt.show()
```

Borough with highest average hourly pickups: Manhattan (1002.285056412618 pickups/hour)



Borough with highest average hourly pickups: Manhattan (1002.285056412618 pickups/hour)



```
In [28]:  # Converting pickup_dt to datetime
df['pickup_dt'] = pd.to_datetime(df['pickup_dt'])

# Adding a column for the hour of the day
df['hour'] = df['pickup_dt'].dt.hour

# Adding a column for the day of the week
df['day_of_week'] = df['pickup_dt'].dt.day_name()

# Number of pickups by hour of the day
hourly_pickups = df.groupby('hour')['pickups'].mean().reset_index()

print(f"Number of pickups by hour of the day:\n{hourly_pickups}")
```

Number of pickups by hour of the day:

	hour	pickups
0	0	333.861411
1	1	276.578638
2	2	205.149573
3	3	153.505181
4	4	138.691581
5	5	148.648718
6	6	192.377500
7	7	237.443994
8	8	275.824247
9	9	292.635922
10	10	292.722449
11	11	291.485950
12	12	287.656250
13	13	284.920816
14	14	292.911184
15	15	304.999182
16	16	310.843393
17	17	315.169381
18	18	336.570728
19	19	346.274128
20	20	351.917275
21	21	353.895935
22	22	359.273242
23	23	359.614064

Number of pickups by hour of the day:

	hour	pickups
0	0	333.861411
1	1	276.578638
2	2	205.149573
3	3	153.505181
4	4	138.691581
5	5	148.648718
6	6	192.377500
7	7	237.443994
8	8	275.824247
9	9	292.635922
10	10	292.722449
11	11	291.485950
12	12	287.656250
13	13	284.920816
14	14	292.911184
15	15	304.999182
16	16	310.843393
17	17	315.169381
18	18	336.570728
19	19	346.274128
20	20	351.917275
21	21	353.895935
22	22	359.273242
23	23	359.614064

```
In [29]: ▶ # Day of the week with the highest number of pickups
day_of_week_pickups = df.groupby('day_of_week')['pickups'].mean().reset_index()
day_with_highest_pickups = day_of_week_pickups.loc[day_of_week_pickups['pickups'] == day_of_week_pickups['pickups'].max()]
print(f"Day of the week with the highest number of pickups:\n{day_with_highest_pickups}")
```

Day of the week with the highest number of pickups:

```
day_of_week    Saturday
pickups        313.247615
Name: 2, dtype: object
```

Day of the week with the highest number of pickups:

```
day_of_week    Saturday
pickups        313.247615
Name: 2, dtype: object
```

2- Weather Impact

```
In [30]: ▶ #Correlation between temperature and the number of pickups
temp_pickup_corr = df['temp'].corr(df['pickups'])

print(f"The correlation between temperature and the number of pickups:\n{temp_pickup_corr}")
```

The correlation between temperature and the number of pickups:

```
0.06813480303893318
```

The correlation between temperature and the number of pickups:
0.0681348030389327

```
In [31]: # Correlation between visibility and the number of pickups  
vsb_pickup_corr = df['vsb'].corr(df['pickups'])  
  
print(f"The Correlation between visibility and the number of pickups:\n
```

The Correlation between visibility and the number of pickups:
-0.0028513311640650245

The Correlation between visibility and the number of pickups:
-0.002851331164065039

```
In [32]: # Correlation between wind speed and the number of pickups  
spd_pickup_corr = df['spd'].corr(df['pickups'])  
  
print(f"The Correlation between wind speed and the number of pickups:\n
```

The Correlation between wind speed and the number of pickups:
-0.00523303497639248

The Correlation between wind speed and the number of pickups:
-0.005233034976392543

```
In [33]: # Correlation between 1-hour precipitation and the number of pickups  
pcp01_pickup_corr = df['pcp01'].corr(df['pickups'])  
  
# Correlation between 6-hour precipitation and the number of pickups  
pcp06_pickup_corr = df['pcp06'].corr(df['pickups'])  
  
# Correlation between 24-hour precipitation and the number of pickups  
pcp24_pickup_corr = df['pcp24'].corr(df['pickups'])  
  
print(f"The Correlation between 1-hour precipitation and the number of  
print(f"The Correlation between 6-hour precipitation and the number of  
print(f"The Correlation between 24-hour precipitation and the number of
```

The Correlation between 1-hour precipitation and the number of pickup
s:
nan

The Correlation between 6-hour precipitation and the number of pickup
s:
nan

The Correlation between 24-hour precipitation and the number of pickup
s:
-0.016615510422116304

The Correlation between 1-hour precipitation and the number of pickups:
nan

The Correlation between 6-hour precipitation and the number of pickups:
nan

The Correlation between 24-hour precipitation and the number of pickups:
-0.01661551042211662

3- Seasonal Trends

```
In [34]: # Convert pickup_dt to datetime
df['pickup_dt'] = pd.to_datetime(df['pickup_dt'])

# Adding a column for the season
def get_season(date):
    if date.month in [12, 1, 2]:
        return 'Winter'
    elif date.month in [3, 4, 5]:
        return 'Spring'
    elif date.month in [6, 7, 8]:
        return 'Summer'
    else:
        return 'Fall'

df['season'] = df['pickup_dt'].apply(get_season)

# Column for the hour of the day
df['hour'] = df['pickup_dt'].dt.hour
```

```
In [35]: # Number of pickups by season
season_pickups = df.groupby('season')['pickups'].mean()

print(f"Number of pickups by season:\n{season_pickups}")
```

```
Number of pickups by season:
season
Fall      274.790931
Spring    286.412525
Summer    299.506847
Winter    266.311539
Name: pickups, dtype: float64
```

```
Number of pickups by season:
season
Fall      274.790931
Spring    286.412525
Summer    299.506847
Winter    266.311539
Name: pickups, dtype: float64
```

```
In [36]: # Average number of pickups during holidays vs non-holidays
df['is_holiday'] = df['hday'] == 'Y'
holiday_pickups = df.groupby('is_holiday')['pickups'].mean()

print(f"Average number of pickups during holidays vs non-holidays:\n{ho
```

```
Average number of pickups during holidays vs non-holidays:
is_holiday
False      282.180652
True       278.659517
Name: pickups, dtype: float64
```

```
Average number of pickups during holidays vs non-holidays:
is_holiday
False      282.180652
True       278.659517
Name: pickups, dtype: float64
```

```
In [37]: # Correlation between snow depth and number of pickups
snow_depth_corr = df['sd'].corr(df['pickups'])

print(f"The correlation between snow depth and number of pickups:\n{sno
```

```
The correlation between snow depth and number of pickups:
-0.025359545201926337
```

```
The correlation between snow depth and number of pickups:
-0.025359545201925526
```

4- Hourly Trends

```
In [38]: # Peak hours for Uber pickups in each borough
peak_hours_borough = df.groupby(['borough', 'hour'])['pickups'].mean().
peak_hours_borough = peak_hours_borough.loc[peak_hours_borough.groupby(

print(f"Peak hours for Uber pickups in each borough:\n{peak_hours_borou
```

```
Peak hours for Uber pickups in each borough:
      borough  hour  pickups
18      Bronx   18    40.496855
46  Brooklyn   22   785.226519
63      EWR    15     0.066298
89  Manhattan   17  1121.000000
118     Queens   22   477.403315
138  Staten Island 18     2.403315
```

Peak hours for Uber pickups in each borough:

	borough	hour	pickups
18	Bronx	18	40.496855
46	Brooklyn	22	785.226519
63	EWB	15	0.066298
89	Manhattan	17	1121.000000
118	Queens	22	477.403315
138	Staten Island	18	2.403315

```
In [39]: # Number of pickups during rush hours (7-9 AM, 5-7 PM)
df['is_rush_hour'] = df['hour'].isin([7, 8, 9, 17, 18, 19])
rush_hour_pickups = df.groupby('is_rush_hour')['pickups'].mean()

print(f"Number of pickups during rush hours (7-9 AM, 5-7 PM):\n{rush_ho
```

```
Number of pickups during rush hours (7-9 AM, 5-7 PM):
is_rush_hour
False      275.735359
True       300.613332
Name: pickups, dtype: float64
```

```
Number of pickups during rush hours (7-9 AM, 5-7 PM):
is_rush_hour
False      275.735359
True       300.613332
Name: pickups, dtype: float64
```

```
In [40]: # Average number of pickups during late-night hours (12 AM - 4 AM)
df['is_late_night'] = df['hour'].isin([0, 1, 2, 3, 4])
late_night_pickups = df.groupby('is_late_night')['pickups'].mean()
print(f"Average number of pickups during late-night hours (12 AM - 4 AM)
```

```
Average number of pickups during late-night hours (12 AM - 4 AM):
is_late_night
False      297.110618
True       222.625892
Name: pickups, dtype: float64
```

```
Average number of pickups during late-night hours (12 AM - 4 AM):
is_late_night
False      297.110618
True       222.625892
Name: pickups, dtype: float64
```

5- Borough Comparison

```
In [41]: # Adding a column for weekends
df['is_weekend'] = df['pickup_dt'].dt.dayofweek >= 5
```



```
In [42]: # Borough comparison during different weather conditions
weather_conditions = ['temp', 'vsb', 'spd', 'pcp01', 'pcp06', 'pcp24',
borough_weather_pickups = df.groupby(['borough'] + weather_conditions)[

print(f"Borough comparison during different weather conditions:\n{borou
```

Borough comparison during different weather conditions:

	borough	temp	vsb	spd	pcp01	pcp06	pcp24	sd
0	Bronx	2.0	10.00	7.0	0.0	0.0	0.090	7.395833
1	Bronx	2.0	10.00	8.0	0.0	0.0	0.090	7.395833
2	Bronx	2.0	10.00	13.0	0.0	0.0	0.090	7.395833
3	Bronx	3.0	10.00	11.0	0.0	0.0	0.090	7.395833
4	Bronx	3.0	10.00	14.0	0.0	0.0	0.090	7.395833
...
17783	Staten Island	88.0	9.10	8.0	0.0	0.0	0.000	0.000000
17784	Staten Island	88.0	10.00	7.0	0.0	0.0	0.000	0.000000
17785	Staten Island	89.0	7.75	3.0	0.0	0.0	0.125	0.000000
17786	Staten Island	89.0	7.75	5.0	0.0	0.0	0.125	0.000000
17787	Staten Island	89.0	9.10	7.0	0.0	0.0	0.000	0.000000

	pickups
0	22.0
1	20.5
2	51.0
3	24.0
4	36.0
...	...
17783	1.0
17784	2.0
17785	3.0
17786	2.0
17787	3.0

[17788 rows x 9 columns]

Borough comparison during different weather conditions:

	borough	temp	vsb	spd	pcp01	pcp06	pcp24	sd	\
0	Bronx	2.0	10.00	7.0	0.0	0.0	0.090	7.395833	
1	Bronx	2.0	10.00	8.0	0.0	0.0	0.090	7.395833	
2	Bronx	2.0	10.00	13.0	0.0	0.0	0.090	7.395833	
3	Bronx	3.0	10.00	11.0	0.0	0.0	0.090	7.395833	
4	Bronx	3.0	10.00	14.0	0.0	0.0	0.090	7.395833	
...	
17783	Staten Island	88.0	9.10	8.0	0.0	0.0	0.000	0.000000	
17784	Staten Island	88.0	10.00	7.0	0.0	0.0	0.000	0.000000	
17785	Staten Island	89.0	7.75	3.0	0.0	0.0	0.125	0.000000	
17786	Staten Island	89.0	7.75	5.0	0.0	0.0	0.125	0.000000	
17787	Staten Island	89.0	9.10	7.0	0.0	0.0	0.000	0.000000	

	pickups
0	22.0
1	20.5
2	51.0
3	24.0
4	36.0
...	...
17783	1.0
17784	2.0
17785	3.0
17786	2.0
17787	3.0

[17788 rows x 9 columns]

```
In [43]: # Borough with the highest increase in pickups during holidays
df['is_holiday'] = df['hday'] == 'Y'
holiday_increase_borough = df.groupby(['borough', 'is_holiday'])['pickups'].max()
holiday_increase_borough = holiday_increase_borough.loc[holiday_increase_borough.ismax()]

print(f"Borough with the highest increase in pickups during holidays:\n
```

Borough with the highest increase in pickups during holidays:

	borough	is_holiday	pickups
0	Bronx	False	30.709296
2	Brooklyn	False	525.110127
5	EWB	True	0.041916
6	Manhattan	False	1002.765565
9	Queens	True	319.945783
10	Staten Island	False	1.606082

Borough with the highest increase in pickups during holidays:

	borough	is_holiday	pickups
0	Bronx	False	30.709296
2	Brooklyn	False	525.110127
5	EWB	True	0.041916
6	Manhattan	False	1002.765565
9	Queens	True	319.945783
10	Staten Island	False	1.606082

```
In [44]: # Number of pickups on weekdays vs weekends for each borough
weekday_weekend_pickups = df.groupby(['borough', 'is_weekend'])['pickups'].max()

print(f"Number of pickups on weekdays vs weekends for each borough:\n{weekday_weekend_pickups}")
```

Number of pickups on weekdays vs weekends for each borough:

	borough	is_weekend	pickups
0	Bronx	False	28.792893
1	Bronx	True	35.098048
2	Brooklyn	False	477.847173
3	Brooklyn	True	640.487179
4	EWB	False	0.025525
5	EWB	True	0.020833
6	Manhattan	False	993.933441
7	Manhattan	True	1022.996795
8	Queens	False	302.674637
9	Queens	True	325.750801
10	Staten Island	False	1.520517
11	Staten Island	True	1.803686

Number of pickups on weekdays vs weekends for each borough:

	borough	is_weekend	pickups
0	Bronx	False	28.792893
1	Bronx	True	35.098048
2	Brooklyn	False	477.847173
3	Brooklyn	True	640.487179
4	EWB	False	0.025525
5	EWB	True	0.020833
6	Manhattan	False	993.933441
7	Manhattan	True	1022.996795
8	Queens	False	302.674637
9	Queens	True	325.750801
10	Staten Island	False	1.520517
11	Staten Island	True	1.803686

6- Weather Extremes

```
In [45]: # Extreme weather conditions effect on pickups
extreme_weather = df[(df['temp'] > 90) | (df['temp'] < 20) | (df['pcp01'] > 0.5) | (df['sd'] > 7.5)]
extreme_weather_pickups = extreme_weather.groupby(['temp', 'pcp01', 'sd']).sum()['pickups']
print(f"Extreme weather conditions effect on pickups:\n{extreme_weather_pickups}")
```

Extreme weather conditions effect on pickups:

temp	pcp01	sd	
0.0	0.0	0.000000	0.000000
2.0	0.0	7.395833	272.300000
3.0	0.0	7.395833	301.000000
4.0	0.0	7.395833	253.788235
5.0	0.0	7.395833	262.515152
...			
55.0	0.0	7.041667	310.500000
56.0	0.0	7.395833	352.833333
57.0	0.0	7.250000	292.428571
58.0	0.0	7.395833	353.000000
59.0	0.0	7.395833	307.285714

Name: pickups, Length: 353, dtype: float64

Extreme weather conditions effect on pickups:

temp	pcp01	sd	
0.0	0.0	0.000000	0.000000
2.0	0.0	7.395833	272.300000
3.0	0.0	7.395833	301.000000
4.0	0.0	7.395833	253.788235
5.0	0.0	7.395833	262.515152
...			
55.0	0.0	7.041667	310.500000
56.0	0.0	7.395833	352.833333
57.0	0.0	7.250000	292.428571
58.0	0.0	7.395833	353.000000
59.0	0.0	7.395833	307.285714

Name: pickups, Length: 353, dtype: float64

```
In [46]: ▶ # Impact of visibility less than 1 mile on pickups
low_visibility = df[df['vsb'] < 1]
low_visibility_pickups = low_visibility['pickups'].mean()
normal_visibility_pickups = df[df['vsb'] >= 1]['pickups'].mean()

print(f"Low Visibility Pickups:\n{low_visibility_pickups}\n")
print(f"Normal Visibility Pickups:\n{normal_visibility_pickups}\n")
```

Low Visibility Pickups:
nan

Normal Visibility Pickups:
282.0452561767637

Low Visibility Pickups:
nan

Normal Visibility Pickups:
282.0452561767637

7- Data Correlations

```
In [47]: ▶ # Correlation between sea level pressure and number of pickups
slp_pickup_corr = df['slp'].corr(df['pickups'])

print(f"The Correlation between sea level pressure and number of pickup
```

The Correlation between sea level pressure and number of pickups:
-0.011690175691148488

The Correlation between sea level pressure and number of pickups:
-0.011690175691148415

```
In [48]: ▶ # Impact of different weather variables on pickups (correlation matrix)
weather_vars = df[['temp', 'dewp', 'spd', 'vsb', 'pickups']]
weather_corr_matrix = weather_vars.corr()
print(f"Impact of different weather variables on pickups (correlation m
```

Impact of different weather variables on pickups (correlation matrix):

	temp	dewp	spd	vsb	pickups
temp	1.000000	0.890263	-0.291404	-0.050096	0.068135
dewp	0.890263	1.000000	-0.320986	-0.316896	0.047266
spd	-0.291404	-0.320986	1.000000	0.122279	-0.005233
vsb	-0.050096	-0.316896	0.122279	1.000000	-0.002851
pickups	0.068135	0.047266	-0.005233	-0.002851	1.000000

Impact of different weather variables on pickups (correlation matrix):

	temp	dewp	spd	vsb	pickups
temp	1.000000	0.890263	-0.291404	-0.050096	0.068135
dewp	0.890263	1.000000	-0.320986	-0.316896	0.047266
spd	-0.291404	-0.320986	1.000000	0.122279	-0.005233
vsb	-0.050096	-0.316896	0.122279	1.000000	-0.002851
pickups	0.068135	0.047266	-0.005233	-0.002851	1.000000

```
In [49]: # Define weather conditions of interest
weather_conditions = ['temp', 'vsb', 'spd', 'pcp01', 'pcp06', 'pcp24',

# Group by holiday status and weather conditions to find the average nu
holiday_weather_pickups = df.groupby([df['hday']=='Y'] + weather_condit

print(f"Average number of pickups on the basis of holiday status and we
```

Average number of pickups on the basis of holiday status and weather c
onditions:

hday	temp	vsb	spd	pcp01	pcp06	pcp24	sd	
False	2.0	10.0	7.000000	0.0	0.0	0.090	7.395833	256.285
714			8.000000	0.0	0.0	0.090	7.395833	267.714
286			13.000000	0.0	0.0	0.090	7.395833	296.333
333	3.0	10.0	11.000000	0.0	0.0	0.090	7.395833	289.428
571			14.000000	0.0	0.0	0.090	7.395833	312.571
429								...
True	83.0	10.0	5.500000	0.0	0.0	0.000	0.000000	386.714
286	84.0	10.0	4.333333	0.0	0.0	0.125	0.000000	365.857
143			6.000000	0.0	0.0	0.125	0.000000	389.428
571			10.000000	0.0	0.0	0.125	0.000000	392.571
429	85.0	10.0	3.666667	0.0	0.0	0.125	0.000000	377.857
143								

Name: pickups, Length: 3259, dtype: float64

Average number of pickups on the basis of holiday status and weather conditions:


	is_holiday	temp	vsb	spd	pcp01	pcp06	pcp24	sd \
0	False	2.0	10.0	7.000000	0.0	0.0	0.090	7.395833
1	False	2.0	10.0	8.000000	0.0	0.0	0.090	7.395833
2	False	2.0	10.0	13.000000	0.0	0.0	0.090	7.395833
3	False	3.0	10.0	11.000000	0.0	0.0	0.090	7.395833
4	False	3.0	10.0	14.000000	0.0	0.0	0.090	7.395833
...
3254	True	83.0	10.0	5.500000	0.0	0.0	0.000	0.000000
3255	True	84.0	10.0	4.333333	0.0	0.0	0.125	0.000000
3256	True	84.0	10.0	6.000000	0.0	0.0	0.125	0.000000
3257	True	84.0	10.0	10.000000	0.0	0.0	0.125	0.000000
3258	True	85.0	10.0	3.666667	0.0	0.0	0.125	0.000000

	pickups
0	256.285714
1	267.714286
2	296.333333
3	289.428571
4	312.571429
...	...
3254	386.714286
3255	365.857143
3256	389.428571
3257	392.571429
3258	377.857143

[3259 rows x 9 columns]

Other than given questions I have tried many things in the EDA to do.

```
In [50]: df.set_index('pickup_dt', inplace=True)
df['pickups'].resample('D').sum().plot(figsize=(15, 6))
plt.title('Daily Uber Pickups')
plt.show()
```



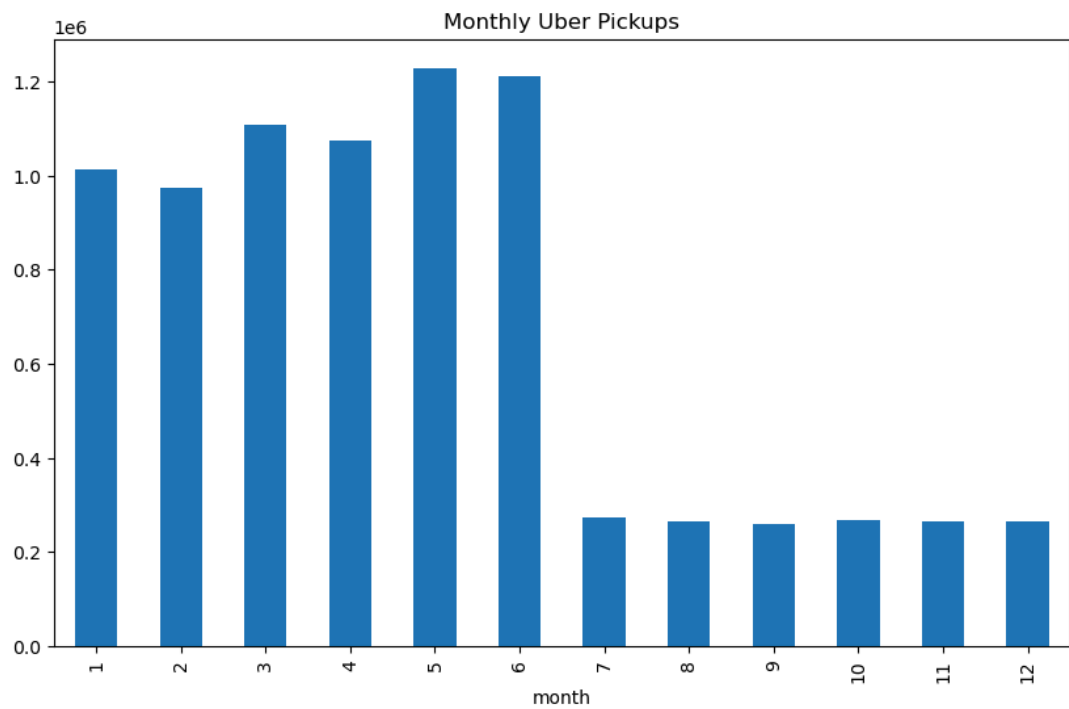
- Time series analysis by which we can see that how the number of pickups are changing across the months

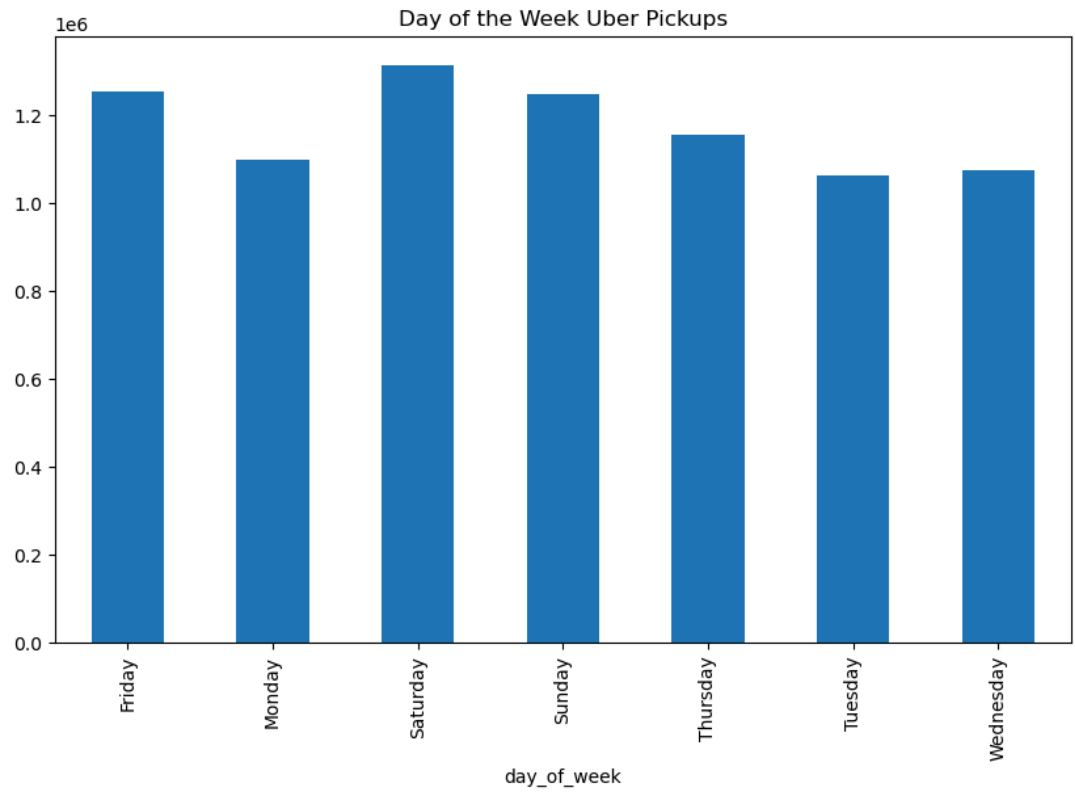
```
In [51]: ▶ df['month'] = df.index.month
df['day_of_week'] = df.index.day_name()

seasonal_pickups = df.groupby('month')['pickups'].sum()
day_of_week_pickups = df.groupby('day_of_week')['pickups'].sum()

seasonal_pickups.plot(kind='bar', figsize=(10, 6))
plt.title('Monthly Uber Pickups')
plt.show()

day_of_week_pickups.plot(kind='bar', figsize=(10, 6))
plt.title('Day of the Week Uber Pickups')
plt.show()
```





In []: ▶